

project_1_starter

September 4, 2018

1 Project 1: Trading with Momentum

1.1 Instructions

Each problem consists of a function to implement and instructions on how to implement the function. The parts of the function that need to be implemented are marked with a `# TODO` comment. After implementing the function, run the cell to test it against the unit tests we've provided. For each problem, we provide one or more unit tests from our `project_tests` package. These unit tests won't tell you if your answer is correct, but will warn you of any major errors. Your code will be checked for the correct solution when you submit it to Udacity.

1.2 Packages

When you implement the functions, you'll only need to use the packages you've used in the classroom, like [Pandas](#) and [Numpy](#). These packages will be imported for you. We recommend you don't add any import statements, otherwise the grader might not be able to run your code.

The other packages that we're importing are `helper`, `project_helper`, and `project_tests`. These are custom packages built to help you solve the problems. The `helper` and `project_helper` module contains utility functions and graph functions. The `project_tests` contains the unit tests for all the problems.

1.2.1 Install Packages

```
In [1]: import sys
        !{sys.executable} -m pip install -r requirements.txt
```

```
Requirement already satisfied: colour==0.1.5 in /opt/conda/lib/python3.6/site-packages (from -r
Collecting cvxpy==1.0.3 (from -r requirements.txt (line 2))
  Downloading https://files.pythonhosted.org/packages/a1/59/2613468ffbbe3a818934d06b81b9f4877fe0
100% || 880kB 647kB/s ta 0:00:01
Requirement already satisfied: cycler==0.10.0 in /opt/conda/lib/python3.6/site-packages/cycler-0
Collecting numpy==1.13.3 (from -r requirements.txt (line 4))
  Downloading https://files.pythonhosted.org/packages/57/a7/e3e6bd9d595125e1abbe162e323fd2d06f6f
100% || 17.0MB 37kB/s eta 0:00:01 20% | | 3.5MB 23.5MB/s eta 0:00:01
Collecting pandas==0.21.1 (from -r requirements.txt (line 5))
  Downloading https://files.pythonhosted.org/packages/3a/e1/6c514df670b887c77838ab856f57783c07e8
100% || 26.2MB 24kB/s eta 0:00:01 8% | | 2.2MB 22.5MB/s eta 0:00:01
Collecting plotly==2.2.3 (from -r requirements.txt (line 6))
```

```

Downloading https://files.pythonhosted.org/packages/99/a6/8214b6564bf4ace9bec8a26e7f89832792be
100% || 1.1MB 601kB/s eta 0:00:01
Requirement already satisfied: pyparsing==2.2.0 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: python-dateutil==2.6.1 in /opt/conda/lib/python3.6/site-packages
Requirement already satisfied: pytz==2017.3 in /opt/conda/lib/python3.6/site-packages (from -r r
Requirement already satisfied: requests==2.18.4 in /opt/conda/lib/python3.6/site-packages (from
Collecting scipy==1.0.0 (from -r requirements.txt (line 11))
Downloading https://files.pythonhosted.org/packages/d8/5e/caa01ba7be11600b6a9d39265440d7b3be3d
100% || 50.0MB 12kB/s eta 0:00:01 11% | 5.9MB 19.2MB/s eta 0
Requirement already satisfied: scikit-learn==0.19.1 in /opt/conda/lib/python3.6/site-packages (f
Requirement already satisfied: six==1.11.0 in /opt/conda/lib/python3.6/site-packages (from -r re
Collecting tqdm==4.19.5 (from -r requirements.txt (line 14))
Downloading https://files.pythonhosted.org/packages/71/3c/341b4fa23cb3abc335207dba057c790f3bb3
100% || 61kB 4.6MB/s eta 0:00:01
Collecting osqp (from cvxpy==1.0.3->-r requirements.txt (line 2))
Downloading https://files.pythonhosted.org/packages/05/42/0ccab82eb6ed0edb83d184928ec864232dc0
100% || 153kB 3.2MB/s eta 0:00:01
Collecting ecos>=2 (from cvxpy==1.0.3->-r requirements.txt (line 2))
Downloading https://files.pythonhosted.org/packages/b6/b4/988b15513b13e8ea2eac65e97d84221ac515
100% || 122kB 4.2MB/s eta 0:00:01
Collecting scs>=1.1.3 (from cvxpy==1.0.3->-r requirements.txt (line 2))
Downloading https://files.pythonhosted.org/packages/b3/fd/6e01c4f4a69fcc6c3db130ba55572089e78e
100% || 143kB 3.3MB/s eta 0:00:01
Collecting multiprocessing (from cvxpy==1.0.3->-r requirements.txt (line 2))
Downloading https://files.pythonhosted.org/packages/7a/ee/b9bf3e171f936743758ef924622d8dd00516
100% || 1.4MB 474kB/s eta 0:00:01 6% | 92kB 9.8MB/s eta 0:00:0
Requirement already satisfied: fastcache in /opt/conda/lib/python3.6/site-packages (from cvxpy==
Requirement already satisfied: toolz in /opt/conda/lib/python3.6/site-packages (from cvxpy==1.0.
Requirement already satisfied: decorator>=4.0.6 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: nbformat>=4.2 in /opt/conda/lib/python3.6/site-packages (from plo
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in /opt/conda/lib/python3.6/site-packages (
Requirement already satisfied: idna<2.7,>=2.5 in /opt/conda/lib/python3.6/site-packages (from re
Requirement already satisfied: urllib3<1.23,>=1.21.1 in /opt/conda/lib/python3.6/site-packages (
Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.6/site-packages (fro
Requirement already satisfied: future in /opt/conda/lib/python3.6/site-packages (from osqp->cvxp
Collecting dill>=0.2.8.1 (from multiprocessing->cvxpy==1.0.3->-r requirements.txt (line 2))
Downloading https://files.pythonhosted.org/packages/6f/78/8b96476f4ae426db71c6e86a8e6a81407f01
100% || 153kB 3.1MB/s eta 0:00:01
Requirement already satisfied: traitlets>=4.1 in /opt/conda/lib/python3.6/site-packages (from nb
Requirement already satisfied: jupyter-core in /opt/conda/lib/python3.6/site-packages (from nbfo
Requirement already satisfied: ipython-genutils in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: jsonschema!=2.5.0,>=2.4 in /opt/conda/lib/python3.6/site-packages
Building wheels for collected packages: cvxpy, plotly, ecos, scs, multiprocessing, dill
Running setup.py bdist_wheel for cvxpy ... done
Stored in directory: /root/.cache/pip/wheels/2b/60/0b/0c2596528665e21d698d6f84a3406c52044c7b4c
Running setup.py bdist_wheel for plotly ... done
Stored in directory: /root/.cache/pip/wheels/98/54/81/dd92d5b0858fac680cd7bdb8800eb26c001dd9f5
Running setup.py bdist_wheel for ecos ... done

```

```

Stored in directory: /root/.cache/pip/wheels/50/91/1b/568de3c087b3399b03d130e71b1fd048ec072c45
Running setup.py bdist_wheel for scs ... done
Stored in directory: /root/.cache/pip/wheels/ff/f0/aa/530ccd478d7d9900b4e9ef5bc5a39e895ce110be
Running setup.py bdist_wheel for multiprocess ... done
Stored in directory: /root/.cache/pip/wheels/8b/36/e5/96614ab62baf927e9bc06889ea794a8e87552b84
Running setup.py bdist_wheel for dill ... done
Stored in directory: /root/.cache/pip/wheels/e2/5d/17/f87cb7751896ac629b435a8696f83ee75b11029f
Successfully built cvxpy plotly ecos scs multiprocess dill
Installing collected packages: numpy, scipy, osqp, ecos, scs, dill, multiprocess, cvxpy, pandas,
  Found existing installation: numpy 1.12.1
    Uninstalling numpy-1.12.1:
      Successfully uninstalled numpy-1.12.1
  Found existing installation: scipy 0.19.1
    Uninstalling scipy-0.19.1:
      Successfully uninstalled scipy-0.19.1
  Found existing installation: dill 0.2.7.1
    Uninstalling dill-0.2.7.1:
      Successfully uninstalled dill-0.2.7.1
  Found existing installation: pandas 0.20.3
    Uninstalling pandas-0.20.3:
      Successfully uninstalled pandas-0.20.3
  Found existing installation: plotly 2.0.15
    Uninstalling plotly-2.0.15:
      Successfully uninstalled plotly-2.0.15
  Found existing installation: tqdm 4.11.2
    Uninstalling tqdm-4.11.2:
      Successfully uninstalled tqdm-4.11.2
Successfully installed cvxpy-1.0.3 dill-0.2.8.2 ecos-2.0.5 multiprocess-0.70.6.1 numpy-1.13.3 os
You are using pip version 9.0.1, however version 18.0 is available.You should consider upgrading

```

1.2.2 Load Packages

```

In [2]: import pandas as pd
        import numpy as np
        import helper
        import project_helper
        import project_tests

```

1.3 Market Data

1.3.1 Load Data

The data we use for most of the projects is end of day data. This contains data for many stocks, but we'll be looking at stocks in the S&P 500. We also made things a little easier to run by narrowing down our range of time period instead of using all of the data.

```

In [3]: df = pd.read_csv('../data/project_1/eod-quotemedia.csv', parse_dates=['date'], index_

```

```
close = df.reset_index().pivot(index='date', columns='ticker', values='adj_close')

print('Loaded Data')
```

Loaded Data

1.3.2 View Data

Run the cell below to see what the data looks like for close.

```
In [4]: project_helper.print_dataframe(close)
```

1.3.3 Stock Example

Let's see what a single stock looks like from the closing prices. For this example and future display examples in this project, we'll use Apple's stock (AAPL). If we tried to graph all the stocks, it would be too much information.

```
In [5]: apple_ticker = 'AAPL'
        project_helper.plot_stock(close[apple_ticker], '{} Stock'.format(apple_ticker))
```

1.4 Resample Adjusted Prices

The trading signal you'll develop in this project does not need to be based on daily prices, for instance, you can use month-end prices to perform trading once a month. To do this, you must first resample the daily adjusted closing prices into monthly buckets, and select the last observation of each month.

Implement the `resample_prices` to resample `close_prices` at the sampling frequency of `freq`.

```
In [53]: def resample_prices(close_prices, freq='M'):
        """
        Resample close prices for each ticker at specified frequency.

        Parameters
        -----
        close_prices : DataFrame
            Close prices for each ticker and date
        freq : str
            What frequency to sample at
            For valid freq choices, see http://pandas.pydata.org/pandas-docs/stable/timeser

        Returns
        -----
        prices_resampled : DataFrame
            Resampled prices for each ticker and date
        """
        # TODO: Implement Function
```

```

        return close_prices.groupby(pd.Grouper(freq=freq)).last()

project_tests.test_resample_prices(resample_prices)

```

Tests Passed

1.4.1 View Data

Let's apply this function to close and view the results.

```

In [54]: monthly_close = resample_prices(close)
        project_helper.plot_resampled_prices(
            monthly_close.loc[:, apple_ticker],
            close.loc[:, apple_ticker],
            '{} Stock - Close Vs Monthly Close'.format(apple_ticker))

```

1.5 Compute Log Returns

Compute log returns (R_t) from prices (P_t) as your primary momentum indicator:

$$R_t = \log_e(P_t) - \log_e(P_{t-1})$$

Implement the `compute_log_returns` function below, such that it accepts a dataframe (like one returned by `resample_prices`), and produces a similar dataframe of log returns. Use Numpy's [log function](#) to help you calculate the log returns.

```

In [55]: def compute_log_returns(prices):
        """
        Compute log returns for each ticker.

        Parameters
        -----
        prices : DataFrame
            Prices for each ticker and date

        Returns
        -----
        log_returns : DataFrame
            Log returns for each ticker and date
        """
        # TODO: Implement Function

        return np.log(prices) - np.log(prices.shift(1))

project_tests.test_compute_log_returns(compute_log_returns)

```

Tests Passed

1.5.1 View Data

Using the same data returned from `resample_prices`, we'll generate the log returns.

```
In [56]: monthly_close_returns = compute_log_returns(monthly_close)
         project_helper.plot_returns(
             monthly_close_returns.loc[:, apple_ticker],
             'Log Returns of {} Stock (Monthly)'.format(apple_ticker))
```

1.6 Shift Returns

Implement the `shift_returns` function to shift the log returns to the previous or future returns in the time series. For example, the parameter `shift_n` is 2 and `returns` is the following:

	Returns				
	A	B	C	D	
2013-07-08	0.015	0.082	0.096	0.020	...
2013-07-09	0.037	0.095	0.027	0.063	...
2013-07-10	0.094	0.001	0.093	0.019	...
2013-07-11	0.092	0.057	0.069	0.087	...
...

the output of the `shift_returns` function would be:

	Shift Returns				
	A	B	C	D	
2013-07-08	NaN	NaN	NaN	NaN	...
2013-07-09	NaN	NaN	NaN	NaN	...
2013-07-10	0.015	0.082	0.096	0.020	...
2013-07-11	0.037	0.095	0.027	0.063	...
...

Using the same `returns` data as above, the `shift_returns` function should generate the following with `shift_n` as -2:

	Shift Returns				
	A	B	C	D	
2013-07-08	0.094	0.001	0.093	0.019	...
2013-07-09	0.092	0.057	0.069	0.087	...
...
...
...	NaN	NaN	NaN	NaN	...
...	NaN	NaN	NaN	NaN	...

Note: The "... " represents data points we're not showing.

```
In [57]: def shift_returns(returns, shift_n):
         """
         Generate shifted returns
```

```

Parameters
-----
returns : DataFrame
    Returns for each ticker and date
shift_n : int
    Number of periods to move, can be positive or negative

Returns
-----
shifted_returns : DataFrame
    Shifted returns for each ticker and date
"""
# TODO: Implement Function

return returns.shift(shift_n)

```

```
project_tests.test_shift_returns(shift_returns)
```

Tests Passed

1.6.1 View Data

Let's get the previous month's and next month's returns.

```

In [58]: prev_returns = shift_returns(monthly_close_returns, 1)
        lookahead_returns = shift_returns(monthly_close_returns, -1)

project_helper.plot_shifted_returns(
    prev_returns.loc[:, apple_ticker],
    monthly_close_returns.loc[:, apple_ticker],
    'Previous Returns of {} Stock'.format(apple_ticker))
project_helper.plot_shifted_returns(
    lookahead_returns.loc[:, apple_ticker],
    monthly_close_returns.loc[:, apple_ticker],
    'Lookahead Returns of {} Stock'.format(apple_ticker))

```

1.7 Generate Trading Signal

A trading signal is a sequence of trading actions, or results that can be used to take trading actions. A common form is to produce a “long” and “short” portfolio of stocks on each date (e.g. end of each month, or whatever frequency you desire to trade at). This signal can be interpreted as rebalancing your portfolio on each of those dates, entering long (“buy”) and short (“sell”) positions as indicated.

Here's a strategy that we will try: > For each month-end observation period, rank the stocks by *previous* returns, from the highest to the lowest. Select the top performing stocks for the long portfolio, and the bottom performing stocks for the short portfolio.

Implement the `get_top_n` function to get the top performing stock for each month. Get the top performing stocks from `prev_returns` by assigning them a value of 1. For all other stocks, give them a value of 0. For example, using the following `prev_returns`:

	Previous Returns						
	A	B	C	D	E	F	G
2013-07-08	0.015	0.082	0.096	0.020	0.075	0.043	0.074
2013-07-09	0.037	0.095	0.027	0.063	0.024	0.086	0.025
...

The function `get_top_n` with `top_n` set to 3 should return the following:

	Previous Returns						
	A	B	C	D	E	F	G
2013-07-08	0	1	1	0	1	0	0
2013-07-09	0	1	0	1	0	1	0
...

Note: You may have to use Panda's `DataFrame.iterrows` with `Series.nlargest` in order to implement the function. This is one of those cases where creating a vectorization solution is too difficult.

```
In [123]: def get_top_n(prev_returns, top_n):
            """
            Select the top performing stocks

            Parameters
            -----
            prev_returns : DataFrame
                Previous shifted returns for each ticker and date
            top_n : int
                The number of top performing stocks to get

            Returns
            -----
            top_stocks : DataFrame
                Top stocks for each ticker and date marked with a 1
            """
            # TODO: Implement Function
            top_stocks = prev_returns.apply(lambda x: x.nlargest(top_n), axis=1)
            top_stocks = top_stocks.applymap(lambda x: 0 if pd.isna(x) else 1)
            top_stocks = top_stocks.astype(np.int64)

            return top_stocks

project_tests.test_get_top_n(get_top_n)
```

Tests Passed

1.7.1 View Data

We want to get the best performing and worst performing stocks. To get the best performing stocks, we'll use the `get_top_n` function. To get the worst performing stocks, we'll also use the `get_top_n` function. However, we pass in `-1*prev_returns` instead of just `prev_returns`. Multiplying by negative one will flip all the positive returns to negative and negative returns to positive. Thus, it will return the worst performing stocks.

```
In [124]: top_bottom_n = 50
          df_long = get_top_n(prev_returns, top_bottom_n)
          df_short = get_top_n(-1*prev_returns, top_bottom_n)
          project_helper.print_top(df_long, 'Longed Stocks')
          project_helper.print_top(df_short, 'Shorted Stocks')
```

10 Most Longed Stocks:

INCY, AMD, AVGO, NFX, NFLX, SWKS, ILMN, NVDA, UAL, MU

10 Most Shorted Stocks:

RRC, CHK, FCX, MRO, FTI, DVN, GPS, WYNN, SPLS, MAT

1.8 Projected Returns

It's now time to check if your trading signal has the potential to become profitable!

We'll start by computing the net returns this portfolio would return. For simplicity, we'll assume every stock gets an equal dollar amount of investment. This makes it easier to compute a portfolio's returns as the simple arithmetic average of the individual stock returns.

Implement the `portfolio_returns` function to compute the expected portfolio returns. Using `df_long` to indicate which stocks to long and `df_short` to indicate which stocks to short, calculate the returns using `lookahead_returns`. To help with calculation, we've provided you with `n_stocks` as the number of stocks we're investing in a single period.

```
In [129]: def portfolio_returns(df_long, df_short, lookahead_returns, n_stocks):
          """
          Compute expected returns for the portfolio, assuming equal investment in each long

          Parameters
          -----
          df_long : DataFrame
              Top stocks for each ticker and date marked with a 1
          df_short : DataFrame
              Bottom stocks for each ticker and date marked with a 1
          lookahead_returns : DataFrame
              Lookahead returns for each ticker and date
          n_stocks: int
              The number number of stocks chosen for each month

          Returns
          -----
          portfolio_returns : DataFrame
```

```

        Expected portfolio returns for each ticker and date
        """
        # TODO: Implement Function

        return (lookahead_returns*(df_long - df_short)) / n_stocks

project_tests.test_portfolio_returns(portfolio_returns)

```

Tests Passed

1.8.1 View Data

Time to see how the portfolio did.

```

In [130]: expected_portfolio_returns = portfolio_returns(df_long, df_short, lookahead_returns, 2
          project_helper.plot_returns(expected_portfolio_returns.T.sum(), 'Portfolio Returns')

```

1.9 Statistical Tests

1.9.1 Annualized Rate of Return

```

In [131]: expected_portfolio_returns_by_date = expected_portfolio_returns.T.sum().dropna()
          portfolio_ret_mean = expected_portfolio_returns_by_date.mean()
          portfolio_ret_ste = expected_portfolio_returns_by_date.sem()
          portfolio_ret_annual_rate = (np.exp(portfolio_ret_mean * 12) - 1) * 100

          print("""
          Mean:                {:.6f}
          Standard Error:      {:.6f}
          Annualized Rate of Return: {:.2f}%
          """).format(portfolio_ret_mean, portfolio_ret_ste, portfolio_ret_annual_rate)

```

```

Mean:                0.003076
Standard Error:      0.002180
Annualized Rate of Return: 3.76%

```

The annualized rate of return allows you to compare the rate of return from this strategy to other quoted rates of return, which are usually quoted on an annual basis.

1.9.2 T-Test

Our null hypothesis (H_0) is that the actual mean return from the signal is zero. We'll perform a one-sample, one-sided t-test on the observed mean return, to see if we can reject H_0 .

We'll need to first compute the t-statistic, and then find its corresponding p-value. The p-value will indicate the probability of observing a mean return equally or more extreme than the one we

observed if the null hypothesis were true. A small p-value means that the chance of observing the mean we observed under the null hypothesis is small, and thus casts doubt on the null hypothesis. It's good practice to set a desired level of significance or alpha (α) *before* computing the p-value, and then reject the null hypothesis if $p < \alpha$.

For this project, we'll use $\alpha = 0.05$, since it's a common value to use.

Implement the `analyze_alpha` function to perform a t-test on the sample of portfolio returns. We've imported the `scipy.stats` module for you to perform the t-test.

Note: `scipy.stats.ttest_1samp` performs a two-sided test, so divide the p-value by 2 to get 1-sided p-value

```
In [133]: from scipy import stats
```

```
def analyze_alpha(expected_portfolio_returns_by_date):
    """
    Perform a t-test with the null hypothesis being that the expected mean return is zero.

    Parameters
    -----
    expected_portfolio_returns_by_date : Pandas Series
        Expected portfolio returns for each date

    Returns
    -----
    t_value
        T-statistic from t-test
    p_value
        Corresponding p-value
    """
    # TODO: Implement Function
    t_test_results = stats.ttest_1samp(expected_portfolio_returns_by_date, 0)
    t_value = t_test_results[0]
    p_value = t_test_results[1] / 2

    return t_value, p_value

project_tests.test_analyze_alpha(analyze_alpha)
```

Tests Passed

1.9.3 View Data

Let's see what values we get with our portfolio. After you run this, make sure to answer the question below.

```
In [134]: t_value, p_value = analyze_alpha(expected_portfolio_returns_by_date)
print("""
Alpha analysis:
t-value:         {:.3f}
```

```
p-value:          {:.6f}
"""".format(t_value, p_value))
```

Alpha analysis:

```
t-value:          1.411
p-value:          0.082517
```

1.9.4 Question: What p-value did you observe? And what does that indicate about your signal?

#TODO: Put Answer In this Cell

I observed a p-value of 0.082517. Since this is greater than our α of 0.05, we must conclude that there is too great a chance of observing an annualized rate of return of 3.76% under our null hypothesis H_0 (that the actual mean return from the signal is zero), and are thus unable to reject this null hypothesis.

In other words, because the p-value is larger than our α value, we can't confidently state that the 3.76% annualized return we observed *was not* due to random chance.

1.10 Submission

Now that you're done with the project, it's time to submit it. Click the submit button in the bottom right. One of our reviewers will give you feedback on your project with a pass or not passed grade. You can continue to the next section while you wait for feedback.